

# Logistic Regression to Perceptron

# Logistic Regression

- (Log) Linear Model – similar to Naïve Bayes
- Doesn't assume features are independent
- Correlated features don't “double count”

# NB vs. LR

- Both compute the dot product
- NB: sum of log probabilities
- LR: logistic function

# NB vs. LR:

## Parameter Learning

- Naïve Bayes:
  - Learn conditional probabilities **independently** by counting
- Logistic Regression:
  - Learn weights **jointly**

# Features for movie review

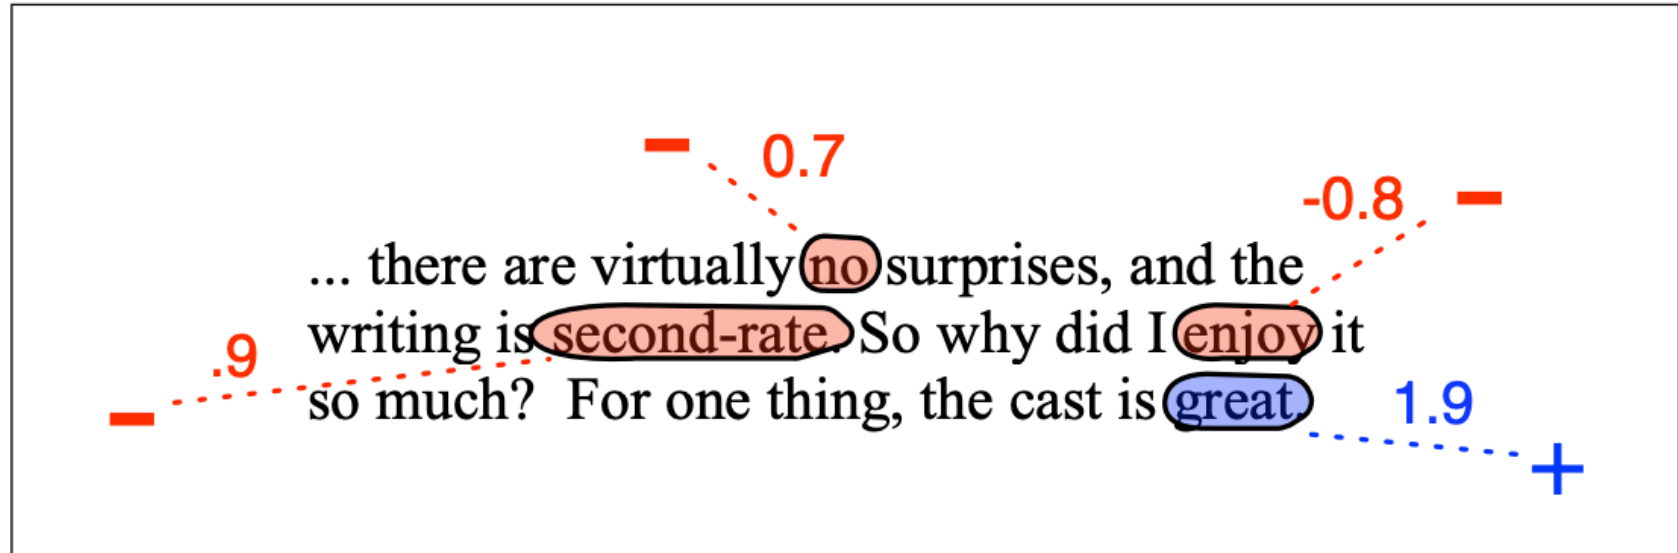
$$f_1(c, x) = \begin{cases} 1 & \text{if "great" } \in x \text{ \& } c = + \\ 0 & \text{otherwise} \end{cases}$$

$$f_2(c, x) = \begin{cases} 1 & \text{if "second-rate" } \in x \text{ \& } c = - \\ 0 & \text{otherwise} \end{cases}$$

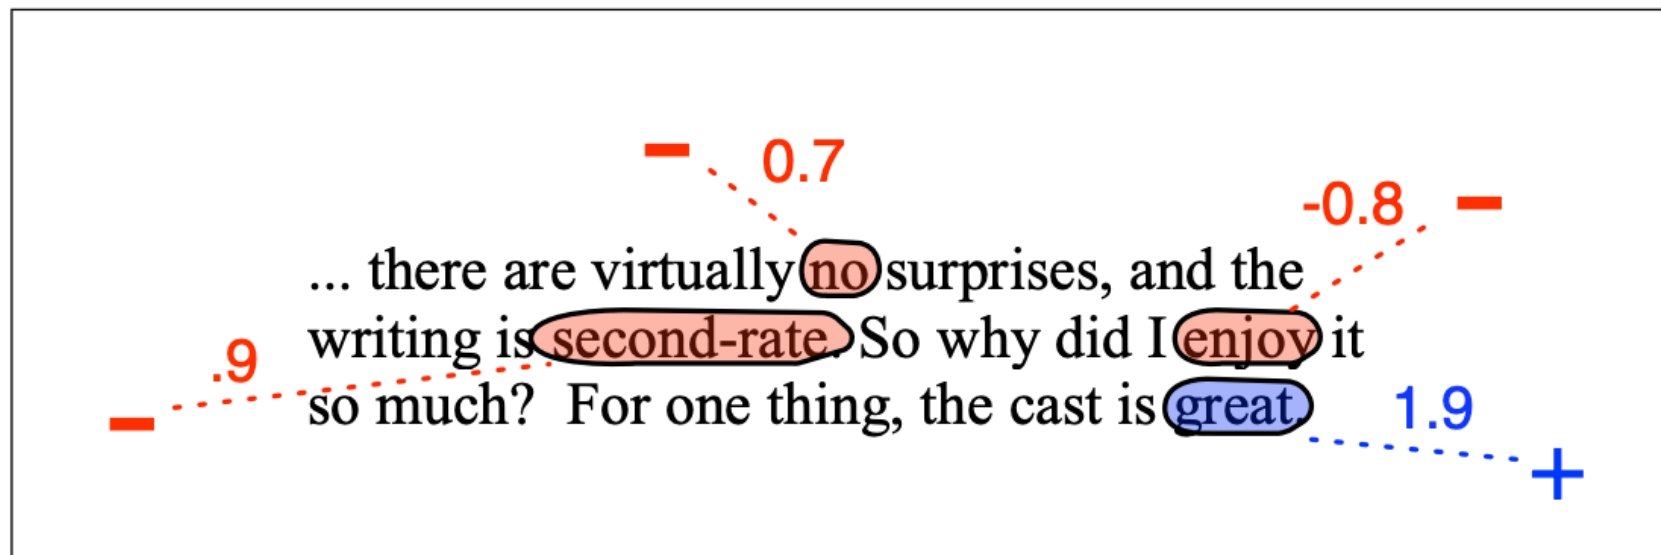
$$f_3(c, x) = \begin{cases} 1 & \text{if "no" } \in x \text{ \& } c = - \\ 0 & \text{otherwise} \end{cases}$$

$$f_4(c, x) = \begin{cases} 1 & \text{if "enjoy" } \in x \text{ \& } c = - \\ 0 & \text{otherwise} \end{cases}$$

# Features for movie review



# Features for movie review



$$P(+|x) = \frac{e^{1.9}}{e^{1.9} + e^{.9}e^{.7}e^{-.8}} = .82$$

$$P(-|x) = \frac{e^{.9}e^{.7}e^{-.8}}{e^{1.9} + e^{.9}e^{.7}e^{-.8}} = .18$$

# Q: what parameters should we choose?

- What is the right value for the weights?
- Maximum Likelihood Principle:
  - Pick the parameters that maximize the probability of the data



# Maximum Likelihood Estimation

- Unfortunately there is no closed form solution
  - (like there was with naïve bayes)
- Solution:
  - Iteratively climb the log-likelihood surface through the derivatives for each weight
- Luckily, the derivatives turn out to be nice

# Gradient ascent

Loop While not converged:


For all features  $j$ , compute and add derivatives

$$w_j^{\text{new}} = w_j^{\text{old}} + \eta \frac{\partial}{\partial w_j} \mathcal{L}(w)$$


$\mathcal{L}(w)$ : Training set log-likelihood

$\left( \frac{\partial \mathcal{L}}{\partial w_1}, \frac{\partial \mathcal{L}}{\partial w_2}, \dots, \frac{\partial \mathcal{L}}{\partial w_n} \right)$  : Gradient vector

# Gradient ascent



# Gradient ascent



# Derivative of Sigmoid

$$\frac{ds(x)}{dx} = \frac{1}{1 + e^{-x}}$$

$$= \left( \frac{1}{1 + e^{-x}} \right)^2 \frac{d}{dx}(1 + e^{-x})$$

$$= \left( \frac{1}{1 + e^{-x}} \right)^2 e^{-x}(-1)$$

$$= \left( \frac{1}{1 + e^{-x}} \right) \left( \frac{1}{1 + e^{-x}} \right) (-e^{-x})$$

$$= \left( \frac{1}{1 + e^{-x}} \right) \left( \frac{-e^{-x}}{1 + e^{-x}} \right)$$

$$= s(x)(1 - s(x))$$

# LR Gradient

$$\frac{\partial \mathcal{L}}{\partial w_j} = \sum_i (y_i - p_i) x_j$$

$j$  -> iterating over features

$i$  -> iterating over training examples

# Perceptron Algorithm

- Algorithm is Very similar to logistic regression
- Not exactly computing gradients

# Perceptron Algorithm

- Algorithm is Very similar to logistic regression
- Not exactly computing gradients

Initialize weight vector  $w = 0$

Loop for  $K$  iterations

    Loop For all training examples  $x_i$

$y_i' = \text{sign}(w * x_i)$

        if  $y_i' \neq y_i$

$w += (y_i - y_i') * x_i$



# Perceptron Notes

- Guaranteed to converge if the data is linearly separable
- Only hyperparameter is maximum number of iterations
- Parameter averaging

# Differences: LR vs. Perceptron

- Batch Learning vs. Online learning
- Perceptron doesn't always make updates

# Online Learning (perceptron)

- Rather than making a full pass through the data, compute gradient and update parameters after each training example.
- Gradients will be less accurate, but the overall effect is to move in the right direction
- Often works well and converges faster than batch learning

# MultiClass Classification

- Q: what if we have more than 2 categories?
  - Sentiment: Positive, Negative, Neutral
  - Document topics: Sports, Politics, Business, Entertainment, ...

# MultiClass Classification

- Q: what if we have more than 2 categories?
  - Sentiment: Positive, Negative, Neutral
  - Document topics: Sports, Politics, Business, Entertainment, ...
- Could train a separate logistic regression model for each category...

# Log-Linear Models

$$P(y|x) \propto e^{w \cdot f(d,y)}$$

$$P(y|x) = \frac{1}{Z(w)} e^{w \cdot f(d,y)}$$

# MultiClass Logistic Regression

$$P(y|x) \propto e^{w \cdot f(d,y)}$$

$$P(y|x) = \frac{1}{Z(w)} e^{w \cdot f(d,y)}$$

$$P(y|x) = \frac{e^{w \cdot f(d,y)}}{\sum_{y' \in Y} e^{w \cdot f(d,y')}}$$

# MultiClass Logistic Regression

- Binary logistic regression:
  - We have one feature vector that matches the size of the vocabulary
- Multiclass in practice:
  - one weight vector for each category

$w_{\text{pos}}$

$w_{\text{neg}}$

$w_{\text{neut}}$



# MultiClass Logistic Regression

- Binary logistic regression:
    - We have one feature vector that matches the size of the output
  - Multi-class logistic regression:
    - One feature vector for each category
- Can represent this in practice with one giant weight vector and repeated features for each category.

$w_{\text{pos}}$

$w_{\text{neg}}$

$w_{\text{neut}}$

# Maximum Likelihood Estimation

$$w_{\text{MLE}} = \operatorname{argmax}_w \log P(y_1, \dots, y_n | x_1, \dots, x_n; w)$$

$$= \operatorname{argmax}_w \sum_i \log P(y_i | x_i; w)$$

$$= \operatorname{argmax}_w \sum_i \log \frac{e^{w \cdot f(x_i, y_i)}}{\sum_{y' \in Y} e^{w \cdot f(x_i, y_i)}}$$

# Multiclass Learning

$$\text{LR : } \frac{\partial \mathcal{L}}{\partial w_j} = \sum_{i=1}^D f_j(y_i, d_i) - \sum_{i=1}^D \sum_{y \in Y} f_j(y, d_i) P(y|d_i)$$

# Multiclass Learning

$$\text{LR : } \frac{\partial \mathcal{L}}{\partial w_j} = \sum_{i=1}^D f_j(y_i, d_i) - \sum_{i=1}^D \sum_{y \in Y} f_j(y, d_i) P(y|d_i)$$

$$\text{Perceptron : } \frac{\partial \mathcal{L}}{\partial w_j} = \sum_{i=1}^D f_j(y_i, d_i) - \sum_{i=1}^D f_j(\arg \max_{y \in Y} P(y|d_i), d_i)$$

# Multiclass Learning

$$\text{LR : } \frac{\partial \mathcal{L}}{\partial w_j} = \sum_{i=1}^D f_j(y_i, d_i) - \sum_{i=1}^D \boxed{\sum_{y \in Y}} f_j(y, d_i) P(y|d_i)$$

$$\text{Perceptron : } \frac{\partial \mathcal{L}}{\partial w_j} = \sum_{i=1}^D f_j(y_i, d_i) - \sum_{i=1}^D f_j(\boxed{\arg \max_{y \in Y}} P(y|d_i), d_i)$$

# MultiClass Perceptron Algorithm

*Initialize weight vector  $w = 0$*

*Loop for  $K$  iterations*

*Loop For all training examples  $x_i$*

$$y_{pred} = \operatorname{argmax}_y (w_y * x_i)$$

$$\text{if } y_{pred} \neq y_i$$

$$w_{y_{gold}} + = x_i$$

$$w_{y_{pred}} - = x_i$$

Q: what if there are only 2 categories?

$$P(y = j|x_i) = \frac{e^{w_j \cdot x_i}}{\sum_k e^{w_k \cdot x_i}}$$

$$P(y = 1|x) = \frac{e^{w_1 \cdot x}}{e^{w_0 \cdot x + w_1 \cdot x - w_1 \cdot x} + e^{w_1 \cdot x}}$$

$$P(y = 1|x) = \frac{e^{w_1 \cdot x}}{e^{w_0 \cdot x - w_1 \cdot x} e^{w_1 \cdot x} + e^{w_1 \cdot x}}$$

$$P(y = 1|x) = \frac{e^{w_1 \cdot x}}{e^{w_1 \cdot x} (e^{w_0 \cdot x - w_1 \cdot x} + 1)}$$

$$P(y = 1|x) = \frac{1}{e^{w_0 \cdot x - w_1 \cdot x} + 1}$$

Q: what if there are only 2 categories?

$$P(y = 1|x) = \frac{1}{e^{-w' \cdot x} + 1}$$



Sigmoid (logistic) function

$$w' = w_1 - w_0$$



# Next

- Regularization
- Markov Model