

Review-3

Perceptron Learning Algorithm

Input: A list T of training examples $\langle \vec{x}_0, y_0 \rangle \dots \langle \vec{x}_n, y_n \rangle$ where

$$\forall i : y_i \in \{+1, -1\}$$

Output: A classifying hyperplane \vec{w}

Randomly initialize \vec{w} ;

while *model \vec{w} makes errors on the training data* **do**

for $\langle \vec{x}_i, y_i \rangle$ *in* T **do**

 Let $\hat{y} = \text{sign}(\vec{w} \cdot \vec{x}_i)$;

if $\hat{y} \neq y_i$ **then**

$$\vec{w} = \vec{w} + y_i \vec{x}_i;$$

end

end

end

Perceptron Learning Algorithm

Input: A list T of training examples $\langle \vec{x}_0, y_0 \rangle \dots \langle \vec{x}_n, y_n \rangle$ where

$$\forall i : y_i \in \{+1, -1\}$$

Output: A classifying hyperplane \vec{w}

Randomly initialize \vec{w} ;

while *model \vec{w} makes errors on the training data* **do**

for $\langle \vec{x}_i, y_i \rangle$ *in* T **do**

 Let $\hat{y} = \text{sign}(\vec{w} \cdot \vec{x}_i)$;

if $\hat{y} \neq y_i$ **then**

$$\vec{w} = \vec{w} + y_i \vec{x}_i;$$

end

end

end

Perceptron Learning Algorithm

Input: A list T of training examples $\langle \vec{x}_0, y_0 \rangle \dots \langle \vec{x}_n, y_n \rangle$ where

$$\forall i : y_i \in \{+1, -1\}$$

Output: A classifying hyperplane \vec{w}

Randomly initialize \vec{w} ;

while *model \vec{w} makes errors on the training data* **do**

for $\langle \vec{x}_i, y_i \rangle$ *in* T **do**

 Let $\hat{y} = \text{sign}(\vec{w} \cdot \vec{x}_i)$;

if $\hat{y} \neq y_i$ **then**

$$\vec{w} = \vec{w} + y_i \vec{x}_i;$$

end

end

end

Perceptron Learning Algorithm

Input: A list T of training examples $\langle \vec{x}_0, y_0 \rangle \dots \langle \vec{x}_n, y_n \rangle$ where

$$\forall i : y_i \in \{+1, -1\}$$

Output: A classifying hyperplane \vec{w}

Randomly initialize \vec{w} ;

while *model \vec{w} makes errors on the training data* **do**

for $\langle \vec{x}_i, y_i \rangle$ *in* T **do**

 Let $\hat{y} = \text{sign}(\vec{w} \cdot \vec{x}_i)$;

if $\hat{y} \neq y_i$ **then**

$$\vec{w} = \vec{w} + y_i \vec{x}_i;$$

end

end

end


*Converges if the training set is
linearly separable*

*May not converge if the training
set is **not** linearly separable*

Support Vector Machines

- A learning method which *explicitly calculates the maximum margin hyperplane* by solving a gigantic quadratic programming minimization problem.
- Among the very highest -performing traditional machine learning techniques.
- But it's relatively slow and quite complicated.


Setting Up the Optimization Problem



The maximum margin can be characterized as a solution to an optimization problem:


$$\begin{aligned} \max. \quad & \frac{2}{\|\vec{w}\|} \\ \text{s.t.} \quad & (\vec{w} \cdot \vec{x} + b) \geq 1, \quad \forall x \text{ of class 1} \\ & (\vec{w} \cdot \vec{x} + b) \leq -1, \quad \forall x \text{ of class 2} \end{aligned}$$

What if it isn't separable?



Project it to someplace where it is!

$$\phi(\langle x, y \rangle) = x^2 + y^2$$



Multiclass Logistic Regression

$$P_w(y|x) = \frac{\exp(w^\top f(x, y))}{\sum_{y' \in \mathcal{Y}} \exp(w^\top f(x, y'))}$$

sum over output
space to normalize

► Compare to binary:

$$P(y = 1|x) = \frac{\exp(w^\top f(x))}{1 + \exp(w^\top f(x))}$$

negative class implicitly had
 $f(x, y=0) = \text{the zero vector}$

Multiclass Logistic Regression

$$P_w(y|x) = \frac{\exp(w^\top f(x, y))}{\sum_{y' \in \mathcal{Y}} \exp(w^\top f(x, y'))}$$

Softmax
function

↗
sum over output
space to normalize

Multiclass Logistic Regression

$$P_w(y|x) = \frac{\exp(w^\top f(x, y))}{\sum_{y' \in \mathcal{Y}} \exp(w^\top f(x, y'))}$$

↗
sum over output
space to normalize

► Training: maximize $\mathcal{L}(x, y) = \sum_{j=1}^n \log P(y_j^* | x_j)$

$$= \sum_{j=1}^n \left(w^\top f(x_j, y_j^*) - \log \sum_y \exp(w^\top f(x_j, y)) \right)$$

Training

► Multiclass logistic regression $P_w(y|x) = \frac{\exp(w^\top f(x, y))}{\sum_{y' \in \mathcal{Y}} \exp(w^\top f(x, y'))}$

► Likelihood $\mathcal{L}(x_j, y_j^*) = w^\top f(x_j, y_j^*) - \log \sum_y \exp(w^\top f(x_j, y))$

$$\frac{\partial}{\partial w_i} \mathcal{L}(x_j, y_j^*) = f_i(x_j, y_j^*) - \frac{\sum_y f_i(x_j, y) \exp(w^\top f(x_j, y))}{\sum_y \exp(w^\top f(x_j, y))}$$

$$\frac{\partial}{\partial w_i} \mathcal{L}(x_j, y_j^*) = f_i(x_j, y_j^*) - \sum_y f_i(x_j, y) P_w(y|x_j)$$

$$\frac{\partial}{\partial w_i} \mathcal{L}(x_j, y_j^*) = \underbrace{f_i(x_j, y_j^*)}_{\text{gold feature value}} - \underbrace{\mathbb{E}_y[f_i(x_j, y)]}_{\text{model's expectation of feature value}}$$

Training

$$\frac{\partial}{\partial w_i} \mathcal{L}(x_j, y_j^*) = f_i(x_j, y_j^*) - \sum_y f_i(x_j, y) P_w(y|x_j)$$

too many drug trials, too few patients

$y^* = \text{Health}$

$$f(x, y = \text{Health}) = [1, 1, 0, 0, 0, 0, 0, 0, 0]$$

$$P_w(y|x) = [0.21, 0.77, 0.02]$$

$$f(x, y = \text{Sports}) = [0, 0, 0, 1, 1, 0, 0, 0, 0]$$

$$\begin{aligned} \text{gradient: } & [1, 1, 0, 0, 0, 0, 0, 0, 0] - 0.21 [1, 1, 0, 0, 0, 0, 0, 0, 0] \\ & - 0.77 [0, 0, 0, 1, 1, 0, 0, 0, 0] - 0.02 [0, 0, 0, 0, 0, 0, 1, 1, 0] \\ & = [0.79, 0.79, 0, -0.77, -0.77, 0, -0.02, -0.02, 0] \end{aligned}$$

update w^\top :

$$\begin{aligned} & [1.3, 0.9, -5, 3.2, -0.1, 0, 1.1, -1.7, -1.3] + [0.79, 0.79, 0, \\ & = [2.09, 1.69, 0, 2.43, -0.87, 0, 1.08, -1.72, 0] \end{aligned}$$

$$P_w(y|x) = \frac{\exp(w^\top f(x, y))}{\sum_{y' \in \mathcal{Y}} \exp(w^\top f(x, y'))}$$

Softmax
function

↪ new $P_w(y|x) = [0.89, 0.10, 0.01]$

Logistic Regression

► Model: $P_w(y|x) = \frac{\exp(w^\top f(x, y))}{\sum_{y' \in \mathcal{Y}} \exp(w^\top f(x, y'))}$

► Inference: $\operatorname{argmax}_y P_w(y|x)$

► Learning: gradient ascent on the discriminative log-likelihood

$$f(x, y^*) - \mathbb{E}_y[f(x, y)] = f(x, y^*) - \sum_y [P_w(y|x) f(x, y)]$$

“towards gold feature value, away from expectation of feature value”

Hidden Markov Models


- ▶ Input $\mathbf{x} = (x_1, \dots, x_n)$ Output $\mathbf{y} = (y_1, \dots, y_n)$
- ▶ Model the sequence of y as a Markov process (dynamics model)
- ▶ Markov property: future is conditionally independent of the past given the present

$$\begin{array}{c} \textcircled{y_1} \longrightarrow \textcircled{y_2} \longrightarrow \textcircled{y_3} \end{array} \quad P(y_3|y_1, y_2) = P(y_3|y_2)$$

- ▶ Lots of mathematical theory about how Markov chains behave
- ▶ If y are tags, this roughly corresponds to assuming that the next tag only depends on the current tag, not anything before

Inference in HMMs

► Input $\mathbf{x} = (x_1, \dots, x_n)$ Output $\mathbf{y} = (y_1, \dots, y_n)$



$$P(\mathbf{y}, \mathbf{x}) = P(y_1) \prod_{i=2}^n P(y_i | y_{i-1}) \prod_{i=1}^n P(x_i | y_i)$$

► Inference problem: $\operatorname{argmax}_{\mathbf{y}} P(\mathbf{y} | \mathbf{x}) = \operatorname{argmax}_{\mathbf{y}} \frac{P(\mathbf{y}, \mathbf{x})}{P(\mathbf{x})}$

Viterbi Algorithm

1. **Initial:** For each state s , calculate

$$\text{score}_1(s) = P(s)P(x_1|s) = \pi_s B_{x_1,s}$$

2. **Recurrence:** For $i = 2$ to n , for every state s , calculate

$$\begin{aligned}\text{score}_i(s) &= \max_{y_{i-1}} P(s|y_{i-1})P(x_i|s)\text{score}_{i-1}(y_{i-1}) \\ &= \max_{y_{i-1}} A_{y_{i-1},s} B_{s,x_i} \text{score}_{i-1}(y_{i-1})\end{aligned}$$

3. **Final state:** calculate

$$\max_{\mathbf{y}} P(\mathbf{y}, \mathbf{x}|\pi, A, B) = \max_s \text{score}_n(s)$$

π : Initial probabilities

A: Transitions

B: Emissions

This only calculates the max. To get final answer (*argmax*),

- keep track of which state corresponds to the max at each step
- build the answer using these back pointers

Example

Sentence: “Learning changes people”

Tagset: NN,VB

				q_F
VB	$9 \cdot 10^{-4}$ (q)	$3.6 \cdot 10^{-7}$ (VB)	$7.2 \cdot 10^{-1}$ 1	$1.44 \cdot 10^{-12}$
NN	$2 \cdot 10^{-4}$ (q)	$10.8 \cdot 10^{-6}$	$7.2 \cdot 10^{-8}$ (VB)	$7.2 \cdot 10^{-9}$ (NN)
Input	Learnin g	change s	People	
Output	VB	VB	NN	

$$P(NN|VB) = 4 \cdot 10^{-1}$$

$$P(NN|NN) = 1 \cdot 10^{-1}$$

$$P(VB|VB) = 1 \cdot 10^{-1}$$

$$P(VB|NN) = 3 \cdot 10^{-1}$$

$$P(VB|q_0) = 3 \cdot 10^{-1}$$

$$P(NN|q_0) = 2 \cdot 10^{-1}$$

$$P(q_f|NN) = 1 \cdot 10^{-1}$$

$$P(q_f|VB) = 2 \cdot 10^{-2}$$

$$P(\text{Learning}|VB) = 3 \cdot 10^{-3}$$

$$P(\text{Learning}|NN) = 1 \cdot 10^{-3}$$

$$P(\text{People}|NN) = 5 \cdot 10^{-2}$$

$$P(\text{People}|VB) = 2 \cdot 10^{-4}$$

$$P(\text{changes}|NN) = 3 \cdot 10^{-3}$$

$$P(\text{changes}|VB) = 4 \cdot 10^{-2}$$

Conditional Random Fields

► HMMs: $P(\mathbf{y}, \mathbf{x}) = P(y_1)P(x_1|y_1)P(y_2|y_1)P(x_2|y_2) \dots$

► CRFs: discriminative models with the following globally-normalized form:


$$P(\mathbf{y}|\mathbf{x}) = \frac{1}{Z} \prod_k \exp(\phi_k(\mathbf{x}, \mathbf{y}))$$

normalizer

↑ any real-valued scoring function of its arguments

K-Nearest Neighbor

- Given training data $\mathcal{D} = \{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_N, \mathbf{y}_N)\}$ and a test point
- Prediction Rule: Look at the K most similar training examples



- For classification: assign the majority class label (**majority voting**)
- For regression: assign the **average response**
- The algorithm requires:
 - Parameter K : number of nearest neighbors to look for
 - **Distance function**: To compute the similarities between examples

K-Nearest Neighbor Algorithm

- Compute the test point's distance from each training point
- Sort the distances in ascending (or descending) order
- Use the sorted distances to select the K nearest neighbors
- Use **majority rule** (for classification) or **averaging** (for regression)

Note: K -Nearest Neighbors is called a *non-parametric* method

- Unlike other supervised learning algorithms, K -Nearest Neighbors doesn't learn an explicit mapping f from the training data
- It simply uses the training data at the test time to make predictions

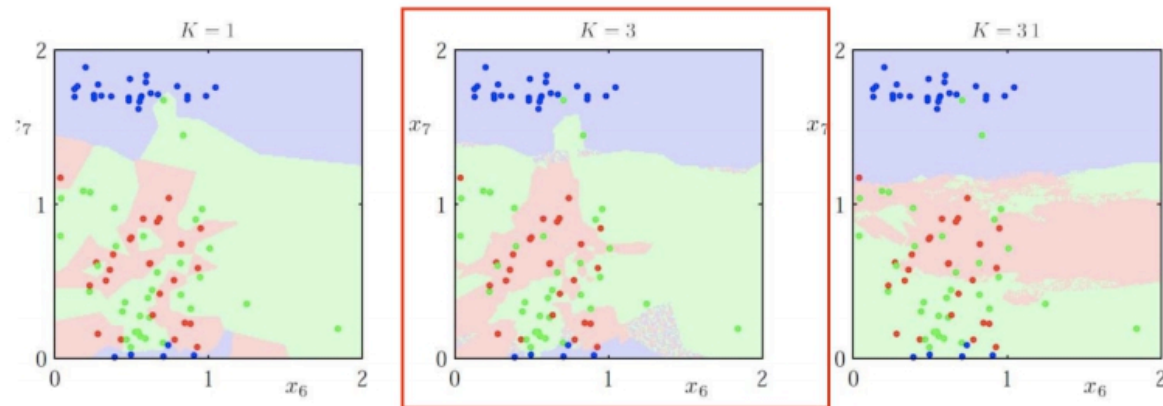
K-NN: Feature Normalization

- Note: Features should be on the same scale
- Example: if one feature has its values in millimeters and another has in centimeters, we would need to normalize
- One way is:
 - Replace x_{im} by $z_{im} = \frac{(x_{im} - \bar{x}_m)}{\sigma_m}$ (make them zero mean, unit variance)
 - $\bar{x}_m = \frac{1}{N} \sum_{i=1}^N x_{im}$: empirical mean of m^{th} feature

K-NN: Other Distance Measure

- Binary-valued features
 - Use Hamming distance: $d(x_i, x_j) = \sum_{m=1}^D \mathbb{I}(x_{im} \neq x_{jm})$
 - Hamming distance counts the number of features where the two examples disagree
- Mixed feature types (some real-valued and some binary-valued)?
 - Can use mixed distance measures
 - E.g., Euclidean for the real part, Hamming for the binary part

K-NN: Choice of K



- Small K
 - Creates many small regions for each class
 - May lead to non-smooth) decision boundaries and overfit
- Large K
 - Creates fewer larger regions
 - Usually leads to smoother decision boundaries (caution: too smooth decision boundary can underfit)
- Choosing K
 - Often data dependent and heuristic based
 - Or using **cross-validation** (using some **held-out data**)
 - In general, a K too small or too big is bad!